

University of Central Florida

Department of Computer Science

COP 3402: Systems Software

Fall 2023

Homework #1 (PM/0-Machine)
This is a team project (1 or 2 people)

See Webcourses for due dates.

I What to Read

Our recommended book is Systems Software: Essential Concepts (by Montagne) in which we recommend reading chapters 1-3.

P-Machine Architecture

The P-machine is a stack machine that conceptually has one memory area called the process address space (PAS). The process address space is divided into two contiguous segments: the “text”, which contains the instructions for the VM to execute and the “stack,” which is organized as a data-stack to be used by the PM/0 CPU.

Registers

The PM/0 has a few built-in registers used for its execution: The registers are named:

- base pointer (BP), which points to the base of the current activation record
- stack pointer (SP), which points to the current top of the stack. The stack grows upwards.,
- program counter (PC), which points to the next instruction to be executed.
- Instruction Register (IR), which store the instruction to be executed

The use of these registers will be explained in detail below. The stack grows upwards.

Instruction Format

The Instruction Set Architecture (ISA) of the PM/0 has instructions that each have three components, which are integers (i.e., they have the C type int) named as follows.

- OP** is the operation code.
- L** indicates the lexicographical level (We will give more details on L below)
- M** depending of the operators it indicates:
 - A number (when OP is LIT or INC).
 - A program address (when OP is JMP, JPC, or CAL).
 - A data address (when OP is LOD, STO)
 - The identity of the arithmetic/relational operation associated to the OPR op-code.
(e.g. OPR 0 2 (ADD) or OPR 0 4 (MUL))

The list of instructions for the ISA can be found in Appendix A and B.

P-Machine Cycles

The PM/0 instruction cycle conceptually does the following for each instruction:

The PM/0 instruction cycle is carried out in two steps. The first step is the fetch cycle, where the instruction pointed to by the program counter (PC) is fetched from the “text” segment, placed in the instruction register (IR) and the PC is incremented to point to the next instruction in the code list. In the second step the instruction in the IR is executed using the “stack” segment. **(This does not mean that the instruction is stored in the “stack segment.”)**

Fetch Cycle:

- 1.- $IR.OP \leftarrow pas[pc]$
 $IR.L \leftarrow pas[pc + 1]$
 $IR.M \leftarrow pas[pc + 2]$
 (note that each instruction need 3 entries in array “TEXT”).
- 2.- $PC \leftarrow PC + 3$.

Execute Cycle:

The op-code (OP) component in the IR register (IR.OP) indicates the operation to be executed. For example, if *IR* encodes the instruction “2 0 2”, then the machine adds the top two elements of the stack, popping them off the stack in the process, and stores the result in the top of the stack (so in the end sp is one less than it was at the start). Note that arithmetic overflows and underflows happen as in C int arithmetic.

PM/0 initial/Default Values

Initial values for PM/0 CPU registers are dependent on the size of the input program, because the “stack” segment in the PAS succeeds the “text” Segment. First you must load the input program into the PAS. Then the values are as follows:

BP = the index immediately following the M value from the last instruction in the program
SP = BP - 1;
PC = 0;

Initial process address space values are all zero:

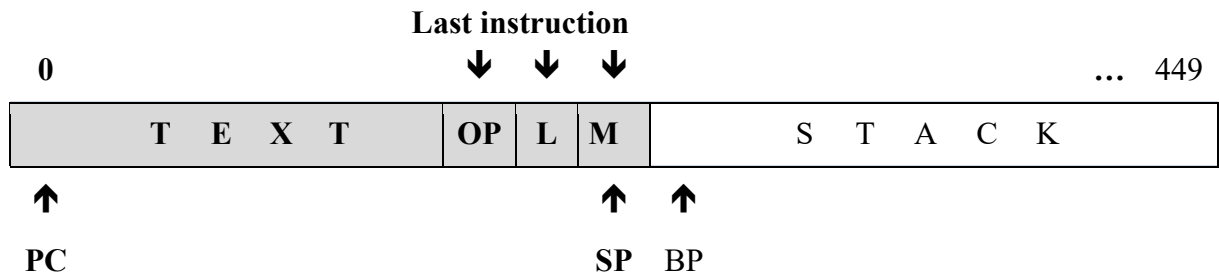
$pas[0] = 0, pas[1] = 0, pas[3] = 0, \dots, [n-1] = 0$.

Constant Values:

ARRAY_SIZE is 512

You will never be given an input file with more than text length greater than 150 lines of code. However, as programs to be run in PM/0 has different length, The values of SP and BP will be set up dynamically once the program had been uploaded.

The figure bellow illustrates the process address space:



Assignment Instructions and Guidelines:

1. The VM must be written in C and must run on Eustis3. If it runs in your PC but not on Eustis, for us it does not run.
2. The input file name should be read as a command line argument at runtime, for example: \$./a.out input.txt (A deduction of 5 points will be applied to submissions that do not implement this).
3. Program output should be printed to the screen, and should follow the formatting of the example in Appendix C. A deduction of 5 points will be applied to submissions that do not implement this.
4. **Submit to Webcourses:**
 - a) A readme text file indicating how to compile and run the VM.
 - b) The source code of your PM/0 VM which should be named "vm.c"
 - c) Student names should be written in the header comment of each source code file, in the readme, and in the comments of the submission
 - d) **Do not change the ISA. Do not add instructions or combine instructions. Do not change the format of the input. If you do so, your grade will be zero.**
 - e) **Include comments in your program. If you do not comments, your grade will be zero.**
 - f) **Do not implement any VM instruction with a function. If you do, a penalty of -100 will be applied to your grade.**
 - g) **On late submissions:**
 - **One day late 10% off.**
 - **Two days late 20% off.**
 - **Submissions will not be accepted after two days.**
 - **Resubmissions are not accepted after two days.**
 - **Your latest submission is the one that will be graded.**

Rubric:

If you submit a program from another semester or we detect plagiarism your grade is F for this course.

Using functions to implement instructions even if only one is implemented that way, means that your grade will be “zero”.

Pointers and handling of dynamic data structures is not allowed. If you do your grade is “zero”

-100 – Does not compile

10 – Compiles

25 – Produces lines of meaningful execution before segfaulting or looping infinitely

5 – Follows IO **specifications** (takes command line argument for input file name and prints output to console)

10 – README.txt containing author names

5 – Fetch cycle is implemented correctly

10 – Well commented source code

5 – Arithmetic instructions are implemented correctly

5 – Read and write instructions are implemented correctly

10 – Load and store instructions are implemented correctly

10 – Call and return instructions are implemented correctly

5 – Follows formatting guidelines correctly, source code is named vm.c

What to submit:

Your program (vm.c) must read in the elf input file (see appendix C) and print out the input file but replacing the opcodes by the mnemonics. For example, instead of printing out 7 0 45, your program must print JMP 0 45. Then print out the executions as indicated in the output file in appendix C.

Please submit:

The vm.c file (your vm implementation)

The input file (ELF)

The output file (execution of the program)

Appendix A

Instruction Set Architecture (ISA) – (eventually we will use “stack” to refer to the stack segment in PAS)

There are 10 arithmetic/logical operations that manipulate the data within the stack segment. These operations are indicated by the **OP** component = 2 (OPR). When an **OPR** instruction is encountered, the **M** component of the instruction is used to select the arithmetic/logical operation to be executed (e.g. to multiply the two elements at the top of the stack, write the instruction “2 0 4”).

ISA:

01 –	LIT 0, M	Pushes a constant value (literal) M onto the stack
02 –	OPR 0, M	Operation to be performed on the data at the top of the stack. (or return from function)
03 –	LOD L, M	Load value to top of stack from the stack location at offset M in AR located L lexicographical levels down
04 –	STO L, M	Store value at top of stack in the stack location at offset M In AR located L lexicographical levels down
05 –	CAL L, M	Call procedure at code index M (generates new Activation Record and $PC \leftarrow M$)
06 –	INC 0, M	Allocate M memory words (increment SP by M).
07 –	JMP 0, M	Jump to instruction M ($PC \leftarrow M$)
08 –	JPC 0, M	Jump to instruction M if top stack element is 0
09 –	SYS 0, 1	Write the top stack element to the screen
	SYS 0, 2	Read in input from the user and store it on top of the stack
	SYS 0, 3	End of program (Set Halt flag to zero)

Appendix B

ISA Pseudo Code

01 – **LIT** 0, M

$sp \leftarrow sp + 1$
 $pas[sp] \leftarrow M;$

02 – **OPR** 0, #

($0 \leq \# \leq 11$)

0 RTN

$sp \leftarrow bp - 1;$
 $bp \leftarrow pas[sp + 2];$
 $pc \leftarrow pas[sp + 3];$

1 ADD

$pas[sp - 1] \leftarrow pas[sp - 1] + pas[sp]$
 $sp \leftarrow sp - 1;$

2 SUB

$pas[sp - 1] \leftarrow pas[sp - 1] - pas[sp]$
 $sp \leftarrow sp - 1;$

3 MUL

$pas[sp - 1] \leftarrow pas[sp - 1] * pas[sp]$
 $sp \leftarrow sp - 1;$

4 DIV

$pas[sp - 1] \leftarrow pas[sp - 1] / pas[sp]$
 $sp \leftarrow sp - 1;$

5 EQL

$pas[sp - 1] \leftarrow pas[sp - 1] == pas[sp]$
 $sp \leftarrow sp - 1;$

6 NEQ

$pas[sp - 1] \leftarrow pas[sp - 1] != pas[sp]$
 $sp \leftarrow sp - 1;$

7 LSS

$pas[sp - 1] \leftarrow pas[sp - 1] < pas[sp]$
 $sp \leftarrow sp - 1;$

8 LEQ

$pas[sp - 1] \leftarrow pas[sp - 1] \leq pas[sp]$
 $sp \leftarrow sp - 1;$

9 GTR

$pas[sp - 1] \leftarrow pas[sp - 1] > pas[sp]$
 $sp \leftarrow sp - 1;$

10 GEQ

$pas[sp - 1] \leftarrow pas[sp - 1] \geq pas[sp]$
 $sp \leftarrow sp - 1;$

03 – LOD L, M	$sp \leftarrow sp + 1;$ $pas[sp] \leftarrow pas[base(bp, L) + M];$
04 – STO L, M	$pas[base(bp, L) + M] \leftarrow pas[sp];$ $sp \leftarrow sp - 1;$
05 - CAL L, M	$pas[sp + 1] \leftarrow base(bp, L);$ /* static link (SL) $pas[sp + 2] \leftarrow bp;$ /* dynamic link (DL) $pas[sp + 3] \leftarrow pc;$ /* return address (RA) $bp \leftarrow sp + 1;$ $pc \leftarrow M;$
06 – INC 0, M	$sp \leftarrow sp + M;$
07 – JMP 0, M	$pc \leftarrow M;$
08 – JPC 0, M	if $pas[sp] == 0$ then { $pc \leftarrow M;$ } $sp \leftarrow sp - 1;$
09 – SYS 0, 1	$printf("%d", pas[sp]);$ $sp \leftarrow sp - 1;$
SYS 0, 2	$sp \leftarrow sp + 1;$ $scanf("%d", pas[sp]);$
SYS 0, 3	Set Halt flag to zero (End of program).

Appendix C

Example of Execution

This example shows how to print the stack after the execution of each instruction.

INPUT FILE (also know as elf file)

For every line, there must be 3 integers representing **OP**, **L** and **M**.

7 0 45

7 0 6

6 0 4

1 0 4

1 0 3

2 0 3

4 1 4

1 0 14

3 1 4

2 0 7

8 0 39

1 0 7

7 0 42

1 0 5

2 0 0

6 0 5

9 0 2

5 0 6

9 0 1

9 0 3

When the input file (program) is read in to be stored in the text segment starting at location 0 in the process address space, each instruction will need three memory locations to be stored. Therefore, the PC must be incremented by 3 in the fetch cycle.

0			3			6			9			12			15			...
7	0	45	7	0	6	6	0	4	1	0	4	1	0	3	2	0	4	etc

The initial CPU register values for the example in this appendix are:

SP = 59;

BP = SP + 1;

PC = 0;

IR.OP = 0, IR.L = 0, IR.M = 0;

Hint: Each instruction uses 3 array elements (in TEXT) and each data value just uses 1 array element (in STACK). Remember Text and Stack coexist in the PAS array

OUTPUT FILE

Print out the execution of the program in the virtual machine, showing the stack and pc, bp, and sp.

NOTE: It is necessary to separate each Activation Record with a bar "|".

			PC	BP	SP	stack
Initial values:			0	60	59	
	JMP	0	45	45	60	59
	INC	0	5	48	60	64 0 0 0 0 0
Please Enter an Integer: 3						
	SYS	0	2	51	60	65 0 0 0 0 0 3
	CAL	0	6	6	66	65 0 0 0 0 0 3
	INC	0	4	9	66	69 0 0 0 0 0 3 60 60 54 0
	LIT	0	4	12	66	70 0 0 0 0 0 3 60 60 54 0 4
	LIT	0	3	15	66	71 0 0 0 0 0 3 60 60 54 0 4 3
	MUL	0	3	18	66	70 0 0 0 0 0 3 60 60 54 0 12
	STO	1	4	21	66	69 0 0 0 0 12 3 60 60 54 0
	LIT	0	14	24	66	70 0 0 0 0 12 3 60 60 54 0 14
	LOD	1	4	27	66	71 0 0 0 0 12 3 60 60 54 0 14 12
	LSS	0	7	30	66	70 0 0 0 0 12 3 60 60 54 0 0
	JPC	0	39	39	66	69 0 0 0 0 12 3 60 60 54 0
	LIT	0	5	42	66	70 0 0 0 0 12 3 60 60 54 0 5
	RTN	0	0	54	60	65 0 0 0 0 12 3
Output result is: 3						
	SYS	0	1	57	60	64 0 0 0 0 12
	SYS	0	3	60	60	64 0 0 0 0 12

Appendix D

Helpful Tips

This function will be helpful to find a variable in a different Activation Record some **L** levels down:

```
/******  
/*      Find base L levels down      */  
/******  
/******
```

```
int base(int BP, int L)  
{  
    int arb = BP;    // arb = activation record base  
    while ( L > 0)    //find base L levels down  
    {  
        arb = pas[arb];  
        L--;  
    }  
    return arb;  
}
```

For example, in the instruction:

STO L, M → You can do stack [base (IR.L) + IR.M]= pas[SP] to store the content of the top of the stack into an AR in the stack, located **L** levels down from the current AR.

Note1: we are working at the CPU level therefore the instruction format must have only 3 fields. Any program whose number of fields in the instruction format is greater than 3 will get a zero.

Note2: If your program does not follow the specifications, your grade will be a zero.

Note3: if any of the instructions is implemented by calling a function, your grade will be zero.

Note4: **Pointers and handling of dynamic data structures is not allowed. If you use them, your grade for HW1 is “zero”.**